

First, let's read in two of the datasets for this week.

```
library(tidyverse)
```

```
key_crop_yields <- read_csv("https://raw.githubusercontent.com/rfordatascience/tidytuesday/master/data/2020/2020-09-01/key_crop_yields.csv")
land_use <- read_csv("https://raw.githubusercontent.com/rfordatascience/tidytuesday/master/data/2020/2020-09-01/land_use_vs_yield_change_in_cereal_production.csv")
```

I'm going to use the `land_use` dataset only to find the top population countries. Let's create a vector of their names.

```
top_countries <- land_use %>%
  janitor::clean_names() %>%
  filter(!is.na(code), entity != "World") %>%
  group_by(entity) %>%
  filter(year == max(year)) %>%
  ungroup() %>%
  slice_max(total_population_gapminder, n = 30) %>%
  pull(entity)
```

```
top_countries
```

```
## [1] "China" "India"
## [3] "United States" "Indonesia"
## [5] "Pakistan" "Brazil"
## [7] "Nigeria" "Bangladesh"
## [9] "Russia" "Mexico"
## [11] "Japan" "Ethiopia"
## [13] "Philippines" "Egypt"
## [15] "Vietnam" "Democratic Republic of Congo"
## [17] "Germany" "Turkey"
## [19] "Iran" "Thailand"
## [21] "United Kingdom" "France"
## [23] "Italy" "South Africa"
## [25] "Tanzania" "Myanmar"
## [27] "Kenya" "South Korea"
## [29] "Colombia" "Spain"
```

Now let's create a tidy version of the crop yields data, for the countries and crops I am interested in.

```
tidy_yields <- key_crop_yields %>%
  janitor::clean_names() %>%
  pivot_longer(wheat_tonnes_per_hectare:bananas_tonnes_per_hectare,
    names_to = "crop", values_to = "yield"
  ) %>%
  mutate(crop = str_remove(crop, "_tonnes_per_hectare")) %>%
  filter(
    crop %in% c("wheat", "rice", "maize", "barley"),
    entity %in% top_countries,
    !is.na(yield)
  )
```

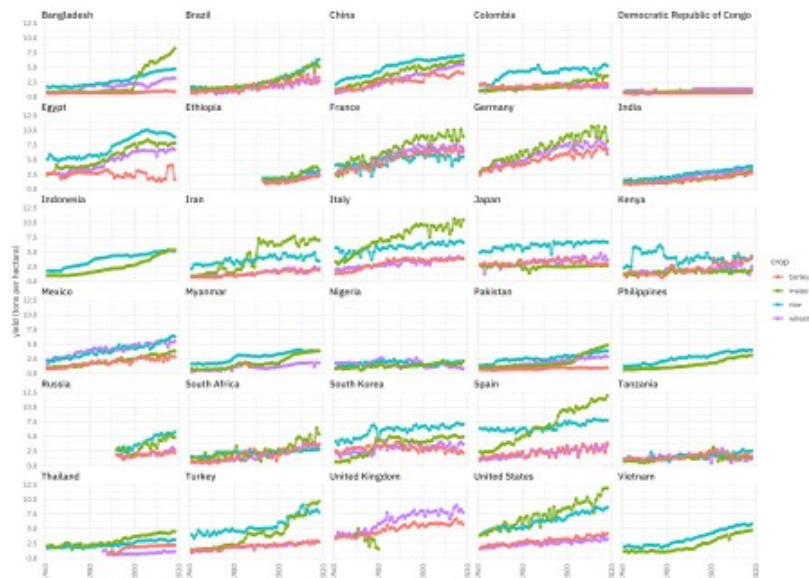
```
tidy_yields
```

```
## # A tibble: 6,032 x 5
##   entity   code  year crop  yield
```

```
##
## 1 Bangladesh BGD      1961 wheat  0.574
## 2 Bangladesh BGD      1961 rice   1.70
## 3 Bangladesh BGD      1961 maize  0.799
## 4 Bangladesh BGD      1961 barley 0.577
## 5 Bangladesh BGD      1962 wheat  0.675
## 6 Bangladesh BGD      1962 rice   1.53
## 7 Bangladesh BGD      1962 maize  0.738
## 8 Bangladesh BGD      1962 barley 0.544
## 9 Bangladesh BGD      1963 wheat  0.607
## 10 Bangladesh BGD     1963 rice   1.77
## # ... with 6,022 more rows
```

This data structure is just right for plotting **crop yield over time**!

```
tidy_yields %>%
  ggplot(aes(year, yield, color = crop)) +
  geom_line(alpha = 0.7, size = 1.5) +
  geom_point() +
  facet_wrap(~entity, ncol = 5) +
  scale_x_continuous(guide = guide_axis(angle = 90)) +
  labs(x = NULL, y = "yield (tons per hectare)")
```



Notice that not all countries produce all crops, but that overall crop yields are *increasing*.

## Many models

Now let's fit a linear model to each country-crop combination.

```
library(tidymodels)

tidy_lm <- tidy_yields %>%
  nest(yields = c(year, yield)) %>%
  mutate(model = map(yields, ~ lm(yield ~ year, data = .x)))

tidy_lm

## # A tibble: 111 x 5
##   entity      code crop  yields      model
##
## 1 Bangladesh BGD  wheat
## 2 Bangladesh BGD  rice
```

```
## 3 Bangladesh BGD  maize
## 4 Bangladesh BGD  barley
## 5 Brazil      BRA  wheat
## 6 Brazil      BRA  rice
## 7 Brazil      BRA  maize
## 8 Brazil      BRA  barley
## 9 China       CHN  wheat
## 10 China      CHN  rice
## # ... with 101 more rows
```

Next, let's tidy() those models to get out the coefficients, and adjust the p-values for multiple comparisons while we're at it.

```
slopes <- tidy_lm %>%
  mutate(coefs = map(model, tidy)) %>%
  unnest(coefs) %>%
  filter(term == "year") %>%
  mutate(p.value = p.adjust(p.value))
```

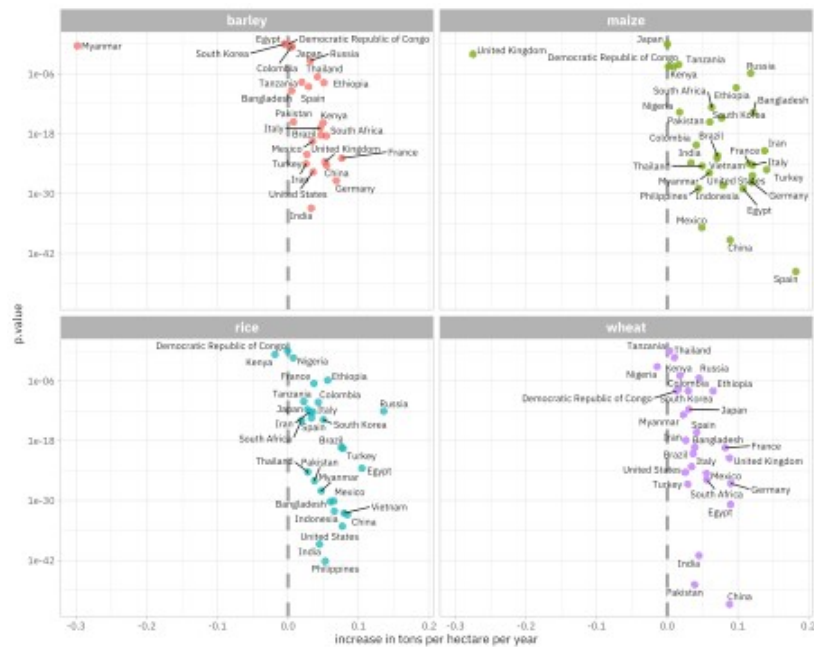
```
slopes
```

```
## # A tibble: 111 x 10
##   entity code crop yields model term estimate std.error statistic
##   <fct> <fct> <fct> <dbl> <dbl> <fct> <dbl> <dbl> <dbl>
##
## 1 Bangla... BGD  wheat   year    0.0389  0.00253    15.4  5.11e-20
## 2 Bangla... BGD  rice    year    0.0600  0.00231    26.0  6.05e-31
## 3 Bangla... BGD  maize   year    0.122   0.0107    11.3  1.82e-14
## 4 Bangla... BGD  barl... year    0.00505  0.000596    8.47  4.34e-10
## 5 Brazil   BRA  wheat   year    0.0366  0.00222    16.5  2.55e-21
## 6 Brazil   BRA  rice    year    0.0755  0.00490    15.4  4.96e-20
## 7 Brazil   BRA  maize   year    0.0709  0.00395    18.0  4.37e-23
## 8 Brazil   BRA  barl... year    0.0466  0.00319    14.6  5.05e-19
## 9 China    CHN  wheat   year    0.0880  0.00141    62.6  1.72e-51
## 10 China   CHN  rice    year    0.0843  0.00289    29.2  1.47e-33
## # ... with 101 more rows
```

## Explore results

Now we can visualize the results of this modeling, which is estimating how crop yields are changing around the world.

```
library(ggplot2)
slopes %>%
  ggplot(aes(estimate, p.value, label = entity)) +
  geom_vline(
    xintercept = 0, lty = 2,
    size = 1.5, alpha = 0.7, color = "gray50"
  ) +
  geom_point(aes(color = crop), alpha = 0.8, size = 2.5, show.legend = FALSE) +
  scale_y_log10() +
  facet_wrap(~crop) +
  geom_text_repel(size = 3, family = "IBMPlexSans") +
  theme_light(base_family = "IBMPlexSans") +
  theme(strip.text = element_text(family = "IBMPlexSans-Bold", size = 12)) +
  labs(x = "increase in tons per hectare per year")
```



- On the x-axis is the slope of these models. Notice that most countries are on the positive side, with increasing crop yields. The further to the right a country is, the larger the increase in crop yield over this time period. Corn yields have increased the most.
- On the y-axis is the p-value, a measure of how surprising the effect we see is under the assumption of no relationship (no change with time). Countries lower in the plots have smaller p-values; we are more certain those are real relationships.

We can extend this to check out how well these models fit the data with `glance()`. This approach for using statistical models to estimate changes in many subgroups at once has been so helpful to me in many situations!